

YZM2031

Data Structures and Algorithms

Week 10: Disjoint Sets

Instructor: Ekrem Çetinkaya

Date: 17.12.2025

Recap

Sorting Algorithms

Elementary Sorts ($O(N^2)$)

- Bubble Sort: Swap adjacent pairs
- Selection Sort: Find min, swap to front
- Insertion Sort: Insert into sorted portion
- Shell Sort: Gap-based insertion sort

Efficient Sorts ($O(N \log N)$)

- Merge Sort: Divide, sort, merge
- Quick Sort: Partition around pivot
- Heap Sort: Build heap, extract max

Key Properties:

- Stability, In-place, Adaptivity

Today's Agenda

1. **Motivation:** Why do we need disjoint sets?
2. **The Equivalence Problem:** Mathematical foundation
3. **Disjoint Set ADT:** Operations and interface
4. **Naive Implementations:** Quick-Find and Quick-Union
5. **Optimizations:** Union by Rank/Size, Path Compression
6. **Applications:** MST, Network Connectivity, Maze Generation

The Grouping Problem

Imagine you're building a social network

You have millions of users. You need to answer questions like:

- Are Alice and Bob in the same friend group?
- How many separate friend groups exist?
- If Alice befriends Charlie, how do we merge their groups?

The Challenge

- Users can be connected through **chains** of friendships
- Groups can **merge dynamically** as new friendships form
- We need **fast** operations on potentially millions of elements
- Cannot afford to traverse entire groups for each query

Real-World Examples

Computer Networks

- Are two computers on the same network?
- Merge networks when a cable is connected
- Count number of isolated networks

Pixels in an Image

- Which pixels belong to the same object?
- Segment image into connected regions
- Count distinct objects in a photo

Electrical Circuits

- Are two components electrically connected?
- What happens when we add a wire?

Road Networks

- Can we drive from city A to city B?
- Which cities are reachable from each other?
- What if we build a new road?

The Core Problem

We need a data structure that can:

1. **Track membership:** Which group does element X belong to?
2. **Merge groups:** Combine two groups into one
3. **Query connectivity:** Are X and Y in the same group?

Performance Requirements

- Elements: Potentially millions
- Operations: Potentially millions
- Each operation should be **nearly constant time**

This is exactly what the Disjoint Set data structure provides

The Equivalence Problem

Equivalence Relations

An equivalence relation \sim on a set S must satisfy three properties:

1. Reflexive

Every element is related to itself: $a \sim a$ for all $a \in S$

2. Symmetric

If a is related to b , then b is related to a : $a \sim b \Rightarrow b \sim a$

3. Transitive

If $a \sim b$ and $b \sim c$, then $a \sim c$

Practice - Identify Equivalence Relations

Which of these are equivalence relations?

1. "Has the same birthday as" on the set of people
2. "Is taller than" on the set of people
3. "Is connected to" in an electrical circuit
4. "Is a sibling of" (including yourself) on the set of people
5. "Lives in the same city as" on the set of people

Answer - Identify Equivalence Relations

1. "Same birthday as" - YES

- Reflexive: You have your own birthday
- Symmetric: If A shares birthday with B, B shares with A
- Transitive: If A and B share, B and C share → A and C share

2. "Is taller than" - NO (not reflexive, not symmetric)

3. "Is connected to" - YES (electrical connectivity)

4. "Is a sibling of" - YES (if we include "sibling of yourself")

5. "Lives in same city as" - YES (same reasoning as birthday)

Equivalence Classes

An equivalence relation **partitions** the set into **disjoint subsets** called **equivalence classes**.

Key Properties

- Every element belongs to **exactly one** equivalence class
- Two elements are in the same class **if and only if** they are related
- Classes are **mutually exclusive** and **collectively exhaustive**

Visualization

Set $S = \{a, b, c, d, e, f, g, h\}$

Equivalence classes: $\{a, b, c\}, \{d, e\}, \{f, g, h\}$



Practice - Building Equivalence Classes

Given: Set $S = \{0, 1, 2, 3, 4, 5, 6, 7\}$

Relations given one by one:

$$1. 0 \sim 1$$

$$2. 2 \sim 3$$

$$3. 4 \sim 5$$

$$4. 6 \sim 7$$

$$5. 0 \sim 2$$

$$6. 4 \sim 6$$

$$7. 0 \sim 4$$

Question: After each relation, how many equivalence classes exist? What are the final classes?

Answer - Building Equivalence Classes

Step	Relation	Classes	Count
Initial	-	{0}, {1}, {2}, {3}, {4}, {5}, {6}, {7}	8
1	$0 \sim 1$	{0,1}, {2}, {3}, {4}, {5}, {6}, {7}	7
2	$2 \sim 3$	{0,1}, {2,3}, {4}, {5}, {6}, {7}	6
3	$4 \sim 5$	{0,1}, {2,3}, {4,5}, {6}, {7}	5
4	$6 \sim 7$	{0,1}, {2,3}, {4,5}, {6,7}	4
5	$0 \sim 2$	{0,1,2,3}, {4,5}, {6,7}	3
6	$4 \sim 6$	{0,1,2,3}, {4,5,6,7}	2
7	$0 \sim 4$	{0,1,2,3,4,5,6,7}	1

The Dynamic Equivalence Problem

The Challenge

- Start with N elements, each in its own equivalence class
- Process a sequence of **union** and **find** operations
- **Union(a, b)**: Merge the classes containing a and b
- **Find(a)**: Return a "representative" of the class containing a

Goal

Answer connectivity queries efficiently:

- Are a and b in the same equivalence class?
- Simply check: `Find(a) == Find(b)`

Disjoint Set ADT

Disjoint Set ADT - Operations

Core Operations

Operation	Description	Returns
<code>makeSet(x)</code>	Create a new set containing only x	<code>void</code>
<code>find(x)</code>	Get the representative of set containing x	representative
<code>union(x, y)</code>	Merge sets containing x and y	<code>void</code>

Derived Operation

Operation	Description	Implementation
<code>connected(x, y)</code>	Are x and y in the same set?	<code>find(x) == find(y)</code>

Key Design Insight

We don't need to enumerate set members

We only need to:

1. **Identify** which set an element belongs to (via representative)
2. **Compare** if two elements are in the same set
3. **Merge** two sets

The Representative

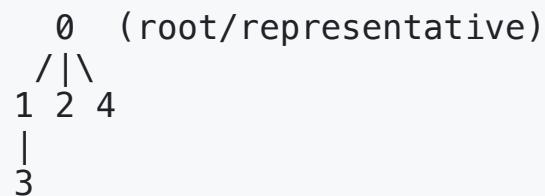
- Each set has exactly **one representative** (also called root or leader)
- All elements in a set have the **same representative**
- Two elements are in the same set \Leftrightarrow same representative

Representation - Trees

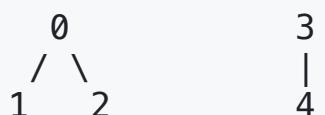
Each set is represented as a tree

- The **root** of the tree is the representative
- Each node points to its **parent**
- The root points to **itself**

Set $\{0, 1, 2, 3, 4\}$



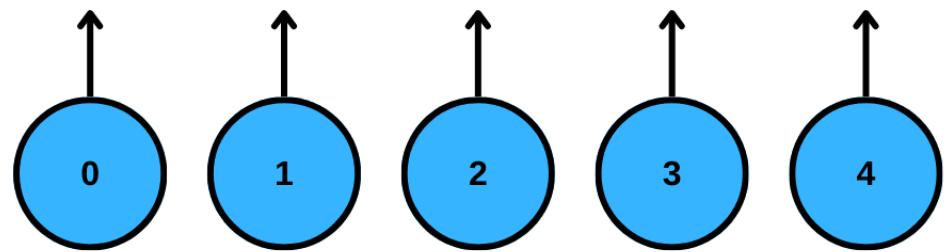
Two sets: $\{0, 1, 2\}$ and $\{3, 4\}$



Array Representation

`s[i] = i if i is a root`

`s[i] = label of the parent if i is not a root`



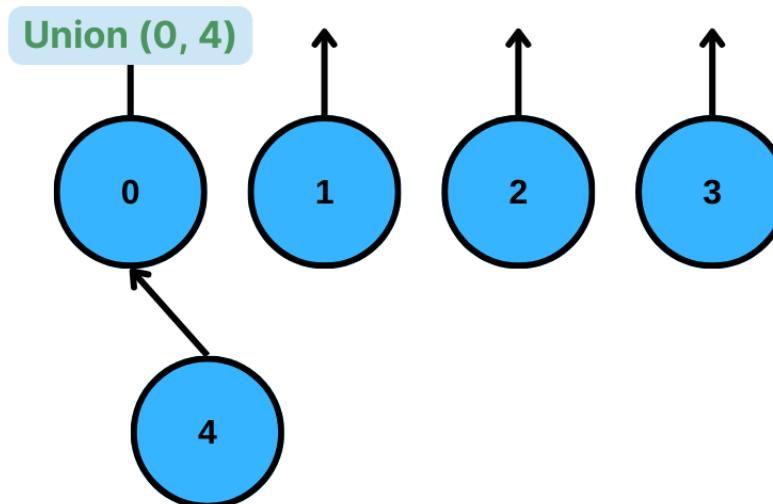
Visualization

Initial State: 5 separate sets, each element is its own representative

```
set = [0, 1, 2, 3, 4]
```

After Union(0, 4): Make 4's root point to 0's root

```
set = [0, 1, 2, 3, 0]
```



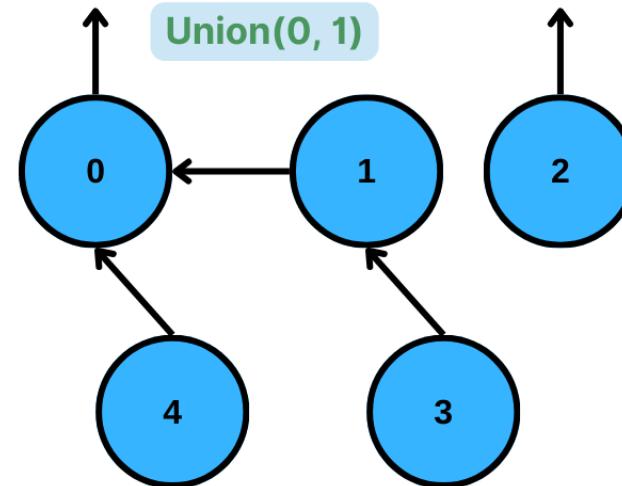
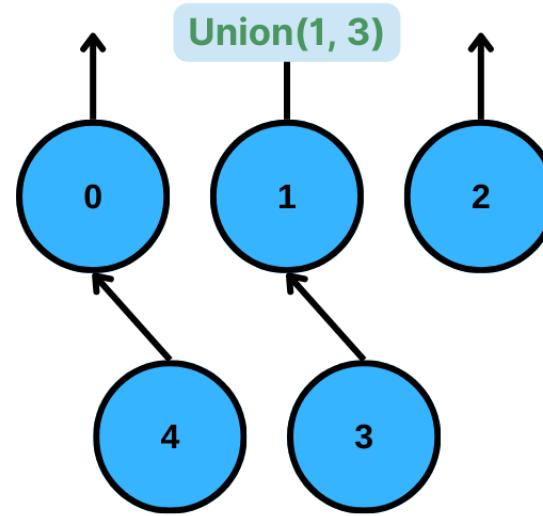
Visualization

After Union (1,3): Make 3's root point to 1's root

```
set = [0, 1, 2, 1, 0]
```

After Union(0, 1): Make 1's root point to 0's root

```
set = [0, 0, 2, 1, 0]
```



Practice - Draw the Forest

Given: 6 elements (0-5), initially separate

Perform these operations (naive quick-union: attach first root under second):

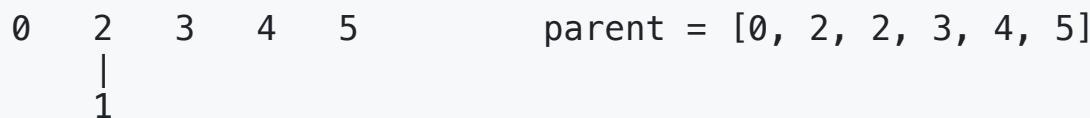
1. `union(1, 2)` - Make root of 1 point to root of 2
2. `union(3, 4)` - Make root of 3 point to root of 4
3. `union(5, 1)` - Make root of 5 point to root of 1
4. `union(3, 5)` - Make root of 3 point to root of 5

Draw the forest and write the parent array after each step.

Answer - Draw the Forest

Initial: 0 1 2 3 4 5 parent = [0, 1, 2, 3, 4, 5]

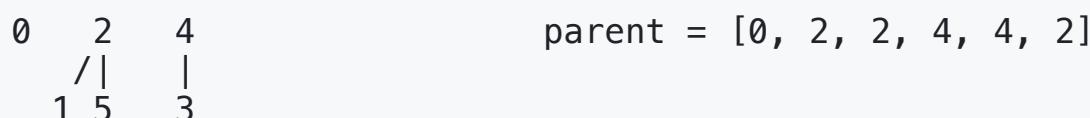
Step 1: union(1,2) – root[1]=1, root[2]=2, parent[1]=2



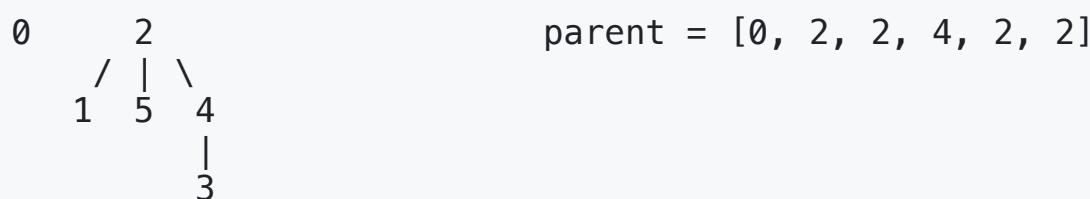
Step 2: union(3,4) – root[3]=3, root[4]=4, parent[3]=4



Step 3: union(5,1) – root[5]=5, root[1]=2, parent[5]=2



Step 4: union(3,5) – root[3]=4, root[5]=2, parent[4]=2



Naive Implementations

Quick-Find (Array-Based Implementation)

Idea

Store the **set ID** (representative) directly for each element.

Index:	0	1	2	3	4	5	6	7
SetID:	[0]	[0]	[0]	[3]	[3]	[5]	[5]	[5]

- Elements 0, 1, 2 are in set "0"
- Elements 3, 4 are in set "3"
- Elements 5, 6, 7 are in set "5"

Key Insight

Every element **directly stores** which set it belongs to.

Quick-Find - Operations

Find(x)

Simply return `id[x]` - $O(1)$

```
int find(int x) {
    return id[x];
}
```

Union(a, b)

Must update **all elements** with `id[a]` to have `id[b]` - $O(N)$

```
void unionSets(int a, int b) {
    int idA = find(a);
    int idB = find(b);
    if (idA == idB) return;

    for (int i = 0; i < n; i++) {
        if (id[i] == idA)
            id[i] = idB;
    }
}
```

Quick-Find - Complete Implementation

```
class QuickFind {
private:
    vector<int> id;

public:
    QuickFind(int n) : id(n) {
        for (int i = 0; i < n; i++)
            id[i] = i; // Initially, each element is its own set
    }

    int find(int x) {
        return id[x]; // O(1) - direct lookup
    }

    void unionSets(int a, int b) {
        int idA = find(a);
        int idB = find(b);
        if (idA == idB) return; // Already in same set

        // Change ALL elements with id[a] to id[b]
        for (int i = 0; i < id.size(); i++) {
            if (id[i] == idA)
                id[i] = idB;
        }
    } // O(N)

    bool connected(int a, int b) {
        return find(a) == find(b); // O(1)
    }
};
```

Practice - Trace Quick-Find

Given: 5 elements (0-4), using Quick-Find

Initial: `id = [0, 1, 2, 3, 4]`

Trace these operations:

1. `union(0, 1)`
2. `union(2, 3)`
3. `union(0, 2)`
4. `connected(1, 3) ?`
5. `connected(3, 4) ?`

Show the `id` array after each union.

Answer - Trace Quick-Find

Initial: id = [0, 1, 2, 3, 4]

Step 1: union(0, 1)

- find(0) = 0, find(1) = 1
- Change all 0s to 1s
- id = [1, 1, 2, 3, 4]

Step 2: union(2, 3)

- find(2) = 2, find(3) = 3
- Change all 2s to 3s
- id = [1, 1, 3, 3, 4]

Step 3: union(0, 2)

- find(0) = 1, find(2) = 3
- Change all 1s to 3s
- id = [3, 3, 3, 3, 4]

Step 4: connected(1, 3)?

- find(1) = 3, find(3) = 3 → TRUE

Step 5: connected(3, 4)?

- find(3) = 3, find(4) = 4 → FALSE

Quick-Find - Analysis

The Problem: Union is too slow!

If we perform N union operations:

```
union(0, 1)    // Scan all N elements
union(0, 2)    // Scan all N elements
union(0, 3)    // Scan all N elements
...
union(0, N-1)  // Scan all N elements
```

Total: $O(N^2)$ for N unions

Complexity Summary

Operation	Time
find	$O(1)$
union	$O(N)$
N unions	$O(N^2)$

Quick-Union (Tree-Based Implementation)

Idea

Store a **parent pointer** for each element. The root points to itself.

Index:	0	1	2	3	4	5	6	7
Parent:	[0]	[0]	[0]	[3]	[3]	[5]	[5]	[5]

Tree visualization:



Key Insight

- Don't store set ID directly
- Store the **path** to the representative (via parent pointers)

Quick-Union - Operations

Find(x)

Follow parent pointers until reaching a root (where `parent[x] == x`)

```
int find(int x) {
    while (x != parent[x])
        x = parent[x];
    return x;
} // O(tree height)
```

Union(a, b)

Make the root of one tree point to the root of the other

```
void unionSets(int a, int b) {
    int rootA = find(a);
    int rootB = find(b);
    if (rootA != rootB)
        parent[rootA] = rootB;
} // O(tree height)
```

Quick-Union - Complete Implementation

```
class QuickUnion {
private:
    vector<int> parent;

public:
    QuickUnion(int n) : parent(n) {
        for (int i = 0; i < n; i++)
            parent[i] = i; // Each element is its own root
    }

    int find(int x) {
        while (x != parent[x]) // Follow chain to root
            x = parent[x];
        return x;
    }

    void unionSets(int a, int b) {
        int rootA = find(a);
        int rootB = find(b);
        if (rootA != rootB)
            parent[rootA] = rootB; // Attach tree A under tree B
    }

    bool connected(int a, int b) {
        return find(a) == find(b);
    }
};
```

Practice - Trace Quick-Union

Given: 5 elements (0-4), using Quick-Union

Initial: `parent = [0, 1, 2, 3, 4]`

Trace these operations:

1. `union(0, 1)` - attach root(0) under root(1)
2. `union(2, 3)` - attach root(2) under root(3)
3. `union(3, 4)` - attach root(3) under root(4)
4. `union(0, 3)` - attach root(0) under root(3)
5. What is `find(0)`? How many steps?

Show the `parent` array and draw trees after each union.

Answer - Trace Quick-Union (1/2)

Initial: parent = [0, 1, 2, 3, 4]
Trees: 0 1 2 3 4

Step 1: union(0, 1) → parent[0] = 1
parent = [1, 1, 2, 3, 4]
Trees: 1 2 3 4
|
0

Step 2: union(2, 3) → parent[2] = 3
parent = [1, 1, 3, 3, 4]
Trees: 1 3 4
| |
0 2

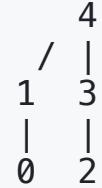
Step 3: union(3, 4) → find(3)=3, find(4)=4, parent[3] = 4
parent = [1, 1, 3, 4, 4]
Trees: 1 4
| |
0 3
|
2

Answer - Trace Quick-Union (2/2)

Step 4: union(0, 3) → find(0)=1, find(3)=4, parent[1] = 4

parent = [1, 4, 3, 4, 4]

Trees:



find(0) = 0→1→4 = 4 (2 steps to reach root)

Quick-Union - The Problem

Trees can become very tall!

Worst case: Always attach larger tree under smaller

union(0, 1): 1
 |
 0

union(1, 2): 2
 |
 1
 |
 0

union(2, 3): 3
 |
 2
 |
 1
 |
 0

After $N - 1$ unions in the worst order: **Linear chain (height = N-1)**

Complexity

Operation	Worst Case
find	$O(N)$
union	$O(N)$
M operations	$O(MN)$

Practice - Worst Case Construction

Question

You have 8 elements (0-7). Give a sequence of 7 union operations that creates the **worst possible tree** (maximum height).

Answer - Worst Case Construction

To create maximum height, always attach to a single-node tree:

```
union(0, 1) → 1-0           h=1
union(1, 2) → 2-1-0         h=2
union(2, 3) → 3-2-1-0       h=3
union(3, 4) → 4-3-2-1-0     h=4
union(4, 5) → 5-4-3-2-1-0   h=5
union(5, 6) → 6-5-4-3-2-1-0 h=6
union(6, 7) → 7-6...0        h=7
```

Now `find(0)` requires 7 steps!

Final Tree:



Comparing Quick-Find vs Quick-Union

Aspect	Quick-Find	Quick-Union
find	$O(1)$	$O(N)$ worst
union	$O(N)$	$O(N)$ worst
Space	$O(N)$	$O(N)$
Advantage	Fast find	Less work per union
Disadvantage	Slow union	Tall trees possible

Neither is good enough!

Both have $O(N^2)$ worst case for N operations.

We need optimizations...

Optimizations

The Key Insight

Why is Quick-Union slow?

Trees can become **tall and skinny** (like linked lists).

Solution Ideas

1. **Union by Rank/Size:** Keep trees **short** by smart merging
2. **Path Compression:** Make trees **flatter** during find operations

Goal

Reduce tree height from $O(N)$ to $O(\log N)$ or even better!

Optimization 1 - Union by Rank

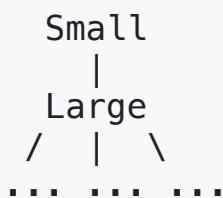
The Problem

Naive union can create tall trees by attaching large trees under small ones.

The Solution

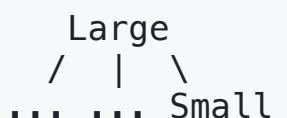
Always attach the **shorter** tree under the **taller** tree.

BAD (height increases):



vs

GOOD (height stays same):



What is Rank?

Rank is an upper bound on the height of a subtree.

Rules

1. Initially, every node has rank 0 (single node, height 0)
2. When unioning two trees:
 - If ranks are **different**: attach smaller rank under larger rank
 - If ranks are **equal**: attach either under the other, **increment** the new root's rank

Why upper bound?

Path compression (later) can reduce actual height, but we don't update rank. So $\text{rank} \geq \text{actual height}$.

Union by Rank - Step by Step

Initial: rank = [0, 0, 0, 0, 0, 0]

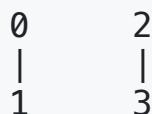
union(0, 1): rank[0]=0, rank[1]=0 (tie!)

- Attach 1 under 0, increment rank[0]
- rank = [1, 0, 0, 0, 0, 0]



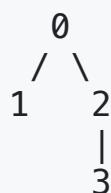
union(2, 3): rank[2]=0, rank[3]=0 (tie!)

- Attach 3 under 2, increment rank[2]
- rank = [1, 0, 1, 0, 0, 0]



union(0, 2): rank[0]=1, rank[2]=1 (tie!)

- Attach 2 under 0, increment rank[0]
- rank = [2, 0, 1, 0, 0, 0]



Union by Rank - Implementation

```
class UnionFindRank {  
private:  
    vector<int> parent;  
    vector<int> rank; // Upper bound on tree height  
  
public:  
    UnionFindRank(int n) : parent(n), rank(n, 0) {  
        for (int i = 0; i < n; i++)  
            parent[i] = i;  
    }  
  
    int find(int x) {  
        while (x != parent[x])  
            x = parent[x];  
        return x;  
    }  
  
    void unionSets(int a, int b) {  
        int rootA = find(a), rootB = find(b);  
        if (rootA == rootB) return;  
  
        if (rank[rootA] < rank[rootB])  
            parent[rootA] = rootB;  
        else if (rank[rootA] > rank[rootB])  
            parent[rootB] = rootA;  
        else {  
            parent[rootB] = rootA;  
            rank[rootA]++;  
        }  
    }  
};
```

Practice - Union by Rank Trace

Given: 8 elements (0-7), all rank 0

Perform these operations with Union by Rank:

1. union(0, 1)
2. union(2, 3)
3. union(4, 5)
4. union(6, 7)
5. union(0, 2)
6. union(4, 6)
7. union(0, 4)

After each step, show the `rank` array and draw the trees.

Answer - Union by Rank Trace (Part 1)

Initial: rank = [0,0,0,0,0,0,0,0], parent = [0,1,2,3,4,5,6,7]

Step 1: union(0,1) – ranks equal, attach 1 under 0, rank[0]++

rank = [1,0,0,0,0,0,0,0]

Trees: 0 2 3 4 5 6 7



Step 2: union(2,3) – ranks equal, attach 3 under 2, rank[2]++

rank = [1,0,1,0,0,0,0,0]

Trees: 0 2 4 5 6 7



Step 3: union(4,5) – ranks equal, attach 5 under 4, rank[4]++

rank = [1,0,1,0,1,0,0,0]

Trees: 0 2 4 6 7



Step 4: union(6,7) – ranks equal, attach 7 under 6, rank[6]++

rank = [1,0,1,0,1,0,1,0]

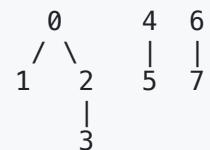
Trees: 0 2 4 6



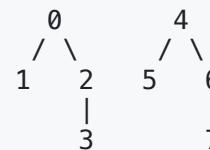
Answer - Union by Rank Trace (Part 2)

After Step 4: rank = [1,0,1,0,1,0,1,0]

Step 5: union(0,2) – rank[0]=1, rank[2]=1 (equal!)
 Attach 2 under 0, rank[0]++
 rank = [2,0,1,0,1,0,1,0]



Step 6: union(4,6) – rank[4]=1, rank[6]=1 (equal!)
 Attach 6 under 4, rank[4]++
 rank = [2,0,1,0,2,0,1,0]



Step 7: union(0,4) – rank[0]=2, rank[4]=2 (equal!)
 Attach 4 under 0, rank[0]++
 rank = [3,0,1,0,2,0,1,0]



Maximum height = 3 (instead of 7 without optimization!)

Union by Rank - Why It Works

Theorem

With Union by Rank, tree height is at most $O(\log N)$.

Proof Sketch

- A tree of rank r has at least 2^r nodes
- Why? When two trees of rank $r - 1$ merge, the result has rank r and at least $2^{r-1} + 2^{r-1} = 2^r$ nodes
- With N nodes: $2^r \leq N \Rightarrow r \leq \log_2 N$

Complexity with Union by Rank

Operation	Time
find	$O(\log N)$
union	$O(\log N)$

Much better than $O(N)!$

Alternative - Union by Size

Instead of rank, track the **number of nodes** in each tree.

```
void unionSets(int a, int b) {
    int rootA = find(a), rootB = find(b);
    if (rootA == rootB) return;

    // Attach smaller tree under larger tree
    if (size[rootA] < size[rootB]) {
        parent[rootA] = rootB;
        size[rootB] += size[rootA];
    } else {
        parent[rootB] = rootA;
        size[rootA] += size[rootB];
    }
}
```

Same $O(\log N)$ guarantee!

Both Union by Rank and Union by Size achieve $O(\log N)$ height.

Optimization 2 - Path Compression

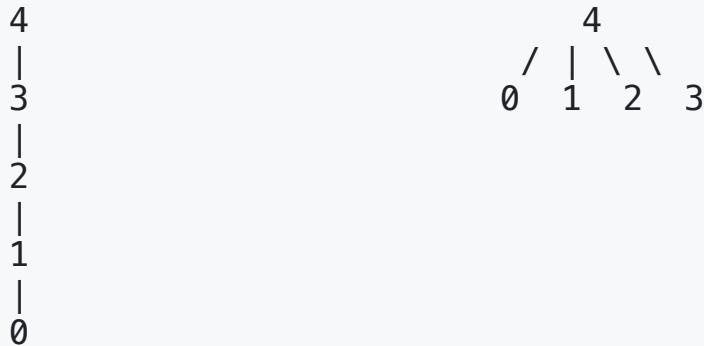
Observation

When we call `find(x)`, we traverse the entire path to the root. But we could **remember** the result

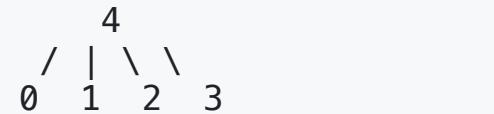
Idea

During `find(x)`, make every node on the path point **directly to the root**.

Before `find(0)`:



After `find(0)` with path compression:



Next time `find(0)`, `find(1)`, `find(2)`, `find(3)` are all $O(1)$!

Path Compression - Recursive Implementation

```
int find(int x) {
    if (x != parent[x])
        parent[x] = find(parent[x]); // Recursively find root, then compress
    return parent[x];
}
```

How it works

1. Recursively find the root
2. On the way back up, update each node to point directly to root
3. Next access is $O(1)$

Path Compression - Iterative Implementation

```
int find(int x) {
    // First pass: find the root
    int root = x;
    while (root != parent[root])
        root = parent[root];

    // Second pass: compress the path
    while (x != root) {
        int next = parent[x];
        parent[x] = root; // Point directly to root
        x = next;
    }
    return root;
}
```

Two-pass algorithm

1. Find the root (standard traversal)
2. Walk path again, updating all nodes to point to root

Practice - Path Compression Trace

Given this tree:

```
6  
|  
5  
|  
4  
|  
3  
|  
2  
|  
1  
|  
0
```

Initial Parent Array:

```
parent = [1, 2, 3, 4, 5, 6, 6]
```

Question

Trace `find(0)` with path compression.

Show the parent array after the operation.

Answer - Path Compression Trace

Initial: parent = [1, 2, 3, 4, 5, 6, 6]

find(0):

Step 1: Find root

0 → 1 → 2 → 3 → 4 → 5 → 6 → 6 (root = 6)

Step 2: Compress path

parent[0] = 6

parent[1] = 6

parent[2] = 6

parent[3] = 6

parent[4] = 6

parent[5] = 6

Final: parent = [6, 6, 6, 6, 6, 6, 6]

Tree after:



Now find(0), find(1), find(2), find(3), find(4), find(5) are all O(1)!

Alternative Compression - Path Splitting

Path Splitting: Make every node point to its **grandparent**.

```
int find(int x) {
    while (x != parent[x]) {
        int next = parent[x];
        parent[x] = parent[parent[x]]; // Point to grandparent
        x = next;
    }
    return x;
}
```

Doesn't compress as aggressively, but simpler (single pass).

Alternative Compression - Path Halving

Path Halving: Make every **other** node point to its grandparent.

```
int find(int x) {
    while (x != parent[x]) {
        parent[x] = parent[parent[x]]; // Point to grandparent
        x = parent[x];              // Skip to grandparent
    }
    return x;
}
```

Even simpler! Also achieves nearly optimal performance.

Combining Both Optimizations

Union by Rank + Path Compression

When we use **both** optimizations together:

Amortized Time Complexity

$O(\alpha(N))$ where α is the inverse Ackermann function.

What is the Ackermann function?

A function that grows **incredibly fast** - faster than any primitive recursive function.

Its inverse $\alpha(N)$ grows **incredibly slowly**.

How Slow is $\alpha(N)$?

N	$\alpha(N)$
1	0
2	1
4	2
16	3
$65536 = 2^{16}$	3
2^{65536} (a number with ~20,000 digits)	4
$2^{2^{65536}}$ (incomprehensibly large)	5

For any practical input size: $\alpha(N) \leq 4$

This means operations are essentially $O(1)$ in practice!

Complete Optimized Implementation

```
class DisjointSet {
private:
    vector<int> parent, rank;

public:
    DisjointSet(int n) : parent(n), rank(n, 0) {
        for (int i = 0; i < n; i++)
            parent[i] = i;
    }

    int find(int x) {
        if (x != parent[x])
            parent[x] = find(parent[x]); // Path compression
        return parent[x];
    }

    void unionSets(int a, int b) {
        int rootA = find(a), rootB = find(b);
        if (rootA == rootB) return;

        // Union by rank
        if (rank[rootA] < rank[rootB]) swap(rootA, rootB);
        parent[rootB] = rootA;
        if (rank[rootA] == rank[rootB]) rank[rootA]++;
    }

    bool connected(int a, int b) { return find(a) == find(b); }
};
```

Complexity Summary

Operation	Quick-Find	Quick-Union	+ Rank	+ Compression
makeSet	$O(1)$	$O(1)$	$O(1)$	$O(1)$
find	$O(1)$	$O(N)$	$O(\log N)$	$O(\alpha(N))$
union	$O(N)$	$O(N)$	$O(\log N)$	$O(\alpha(N))$
M ops	$O(MN)$	$O(MN)$	$O(M \log N)$	$O(M\alpha(N))$

Key Takeaway

With both optimizations, M operations on N elements takes $O(M \cdot \alpha(N)) \approx O(M)$

This is essentially optimal!

Practice - Full Trace with Both Optimizations

Given: 6 elements, using Union by Rank + Path Compression

Operations:

1. union(0, 1)
2. union(2, 3)
3. union(0, 2)
4. union(4, 5)
5. union(0, 4)
6. find(5) (show path compression)

Trace parent and rank arrays. Draw trees.

Answer - Full Trace (1/3)

```
Initial: parent = [0,1,2,3,4,5], rank = [0,0,0,0,0,0]
```

```
union(0,1): rank[0]=rank[1]=0, attach 1 under 0, rank[0]++
parent = [0,0,2,3,4,5], rank = [1,0,0,0,0,0]
```

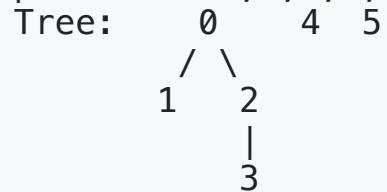
Tree: 0 2 3 4 5
|
1

```
union(2,3): rank[2]=rank[3]=0, attach 3 under 2, rank[2]++
parent = [0,0,2,2,4,5], rank = [1,0,1,0,0,0]
```

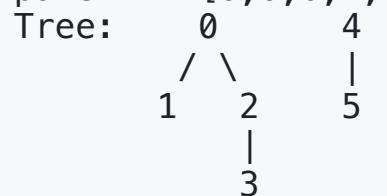
Tree: 0 2 4 5
| |
1 3

Answer - Full Trace (2/3)

```
union(0,2): find(0)=0, find(2)=2, rank[0]=rank[2]=1, attach 2 under 0
parent = [0,0,0,2,4,5], rank = [2,0,1,0,0,0]
```



```
union(4,5): rank[4]=rank[5]=0, attach 5 under 4, rank[4]++
parent = [0,0,0,2,4,4], rank = [2,0,1,0,1,0]
```



Answer - Full Trace (3/3)

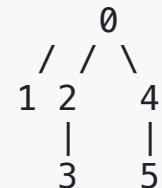
After union(4,5): parent = [0,0,0,2,4,4], rank = [2,0,1,0,1,0]

union(0,4): find(0)=0, find(4)=4, rank[0]=2 > rank[4]=1

Attach 4 under 0

parent = [0,0,0,2,0,4], rank = [2,0,1,0,1,0]

Tree:



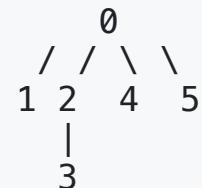
find(5) with path compression:

5 → 4 → 0 (root = 0)

Compress: parent[5] = 0, parent[4] = 0 (already is)

parent = [0,0,0,2,0,0], rank = [2,0,1,0,1,0]

Tree after:



Now find(5) is O(1)!

Applications

Application 1 - Kruskal's MST Algorithm

Problem: Minimum Spanning Tree

Given a weighted, connected graph, find edges that:

1. Connect **all vertices**
2. Have **minimum total weight**
3. Form **no cycles** (exactly $V - 1$ edges for V vertices)

Why Union-Find?

We need to efficiently answer: "Would adding this edge create a cycle?"

An edge (u, v) creates a cycle if u and v are **already connected**.

Kruskal's Algorithm

Algorithm

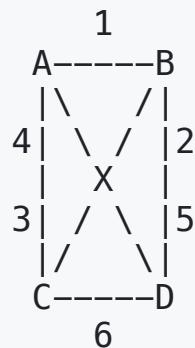
1. Sort all edges by weight (ascending)
2. Initialize Union-Find with V vertices
3. For each edge (u, v, weight) in sorted order:
 - a. If $\text{find}(u) \neq \text{find}(v)$: // u and v in different components
 - Add edge to MST
 - $\text{union}(u, v)$
 - b. Else:
 - Skip edge (would create cycle)
4. Stop when MST has $V-1$ edges

Time Complexity

- Sorting: $O(E \log E)$
- Union-Find operations: $O(E \cdot \alpha(V))$
- Total: $O(E \log E)$

Practice - Kruskal's MST

Given this weighted graph:



Edges: A-B(1), A-C(4), A-D(3), B-C(2), B-D(5), C-D(6)

Trace Kruskal's algorithm step by step.

Answer - Kruskal's MST

Sorted edges: A-B(1), B-C(2), A-D(3), A-C(4), B-D(5), C-D(6)

Initial: {A}, {B}, {C}, {D}

Edge A-B(1): find(A)≠find(B), ADD, union(A,B)

Sets: {A,B}, {C}, {D}

MST: [A-B]

Edge B-C(2): find(B)=A, find(C)=C, ADD, union(B,C)

Sets: {A,B,C}, {D}

MST: [A-B, B-C]

Edge A-D(3): find(A)=A, find(D)=D, ADD, union(A,D)

Sets: {A,B,C,D}

MST: [A-B, B-C, A-D]

MST complete! (3 edges for 4 vertices)

Edge A-C(4): find(A)=find(C)=A, SKIP (would create cycle)

Edge B-D(5): find(B)=find(D)=A, SKIP

Edge C-D(6): find(C)=find(D)=A, SKIP

MST weight: 1 + 2 + 3 = 6

Kruskal's - Implementation

```
struct Edge {
    int u, v, weight;
    bool operator<(const Edge& e) const { return weight < e.weight; }
};

int kruskalMST(int n, vector<Edge>& edges) {
    sort(edges.begin(), edges.end()); // Sort by weight
    DisjointSet ds(n);

    int mstWeight = 0;
    int edgesAdded = 0;

    for (const Edge& e : edges) {
        if (ds.find(e.u) != ds.find(e.v)) {
            ds.unionSets(e.u, e.v);
            mstWeight += e.weight;
            edgesAdded++;

            if (edgesAdded == n - 1) break; // MST complete
        }
    }

    return mstWeight;
}
```

Application 2 - Cycle Detection

Problem

Given an undirected graph, determine if it contains a cycle.

Key Insight

A cycle exists if and only if we try to add an edge between two **already-connected** vertices.

Algorithm

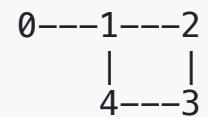
1. Initialize Union-Find with V vertices
2. For each edge (u, v) :
 - a. If $\text{find}(u) == \text{find}(v)$:
 - CYCLE FOUND! (u and v already connected)
 - b. Else:
 - $\text{union}(u, v)$
3. If no cycle found, graph is acyclic (a forest)

Practice - Cycle Detection

Given edges in order: (0,1), (1,2), (2,3), (3,4), (4,1)

Trace the algorithm. Does a cycle exist?

Graph being built:



Answer - Cycle Detection

Initial: $\{0\}, \{1\}, \{2\}, \{3\}, \{4\}$

Edge (0,1): $\text{find}(0)=0, \text{find}(1)=1$, different $\rightarrow \text{union}(0,1)$
Sets: $\{0,1\}, \{2\}, \{3\}, \{4\}$

Edge (1,2): $\text{find}(1)=0, \text{find}(2)=2$, different $\rightarrow \text{union}(1,2)$
Sets: $\{0,1,2\}, \{3\}, \{4\}$

Edge (2,3): $\text{find}(2)=0, \text{find}(3)=3$, different $\rightarrow \text{union}(2,3)$
Sets: $\{0,1,2,3\}, \{4\}$

Edge (3,4): $\text{find}(3)=0, \text{find}(4)=4$, different $\rightarrow \text{union}(3,4)$
Sets: $\{0,1,2,3,4\}$

Edge (4,1): $\text{find}(4)=0, \text{find}(1)=0$, SAME!
 \rightarrow CYCLE DETECTED!

The cycle is: 1-2-3-4-1

Application 3 - Dynamic Connectivity

Problem

Support a stream of operations:

- `connect(a, b)` : Add an edge between a and b
- `query(a, b)` : Are a and b connected?

Solution

```
class DynamicConnectivity {
    DisjointSet ds;
public:
    DynamicConnectivity(int n) : ds(n) {}

    void connect(int a, int b) {
        ds.unionSets(a, b);
    }

    bool query(int a, int b) {
        return ds.find(a) == ds.find(b);
    }
};
```

Each operation is $O(\alpha(N))$ - essentially constant!

Application 4 - Counting Connected Components

Problem

Given n nodes and a list of edges, count distinct connected components.

Solution

```
int countComponents(int n, vector<pair<int,int>>& edges) {
    DisjointSet ds(n);

    for (auto& [u, v] : edges) {
        ds.unionSets(u, v);
    }

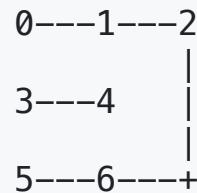
    // Count distinct roots
    unordered_set<int> roots;
    for (int i = 0; i < n; i++) {
        roots.insert(ds.find(i));
    }
    return roots.size();
}
```

Alternative: Track component count during unions (decrement when merging).

Practice - Count Components

Given: $n = 7$, edges = $[(0,1), (1,2), (3,4), (5,6), (2,6)]$

How many connected components exist?



Answer - Count Components

Initial: 7 components (each node is its own)

Edge (0,1): merge {0} and {1} → 6 components
Edge (1,2): merge {0,1} and {2} → 5 components
Edge (3,4): merge {3} and {4} → 4 components
Edge (5,6): merge {5} and {6} → 3 components
Edge (2,6): find(2) in {0,1,2}, find(6) in {5,6}
 merge → 2 components

Final components:

{0, 1, 2, 5, 6}
{3, 4}

Answer: 2 components

Application 5 - Image Segmentation

Connected Component Labeling

Binary Image:	Component Labels:
1 1 0 0 1	A A 0 0 B
1 1 0 1 1	A A 0 B B
0 0 0 1 0	0 0 0 B 0
1 1 1 0 0	C C C 0 0

Algorithm

1. Create Union-Find with one element per pixel
2. Scan image, union adjacent "1" pixels
3. Each root represents a distinct component

Summary

Disjoint Sets (Union-Find)

- Purpose: Track dynamic equivalence / connectivity
- Operations: `makeSet`, `find`, `union`

Optimizations

Technique	Effect
Union by Rank/Size	Keeps trees balanced $\rightarrow O(\log N)$
Path Compression	Flattens trees \rightarrow nearly $O(1)$
Combined	$O(\alpha(N))$ per operation

Key Point

$\alpha(N) \leq 4$ for all practical inputs \rightarrow essentially constant time!

Applications Summary

Application	How Union-Find Helps
Kruskal's MST	Detect if edge creates cycle
Cycle Detection	Check if endpoints already connected
Dynamic Connectivity	Answer "are X and Y connected?"
Connected Components	Count distinct groups
Image Segmentation	Group adjacent similar pixels
Network Connectivity	Track which computers can communicate

Key Takeaways

1. **Simple idea:** Trees to represent sets, roots as representatives
2. **Naive is slow:** Both Quick-Find and Quick-Union are $O(N)$
3. **Optimizations matter:** Rank + Compression gives $O(\alpha(N))$
4. **Practically constant:** $\alpha(N) \leq 4$ for any realistic input
5. **Wide applications:** MST, cycles, connectivity, segmentation
6. **Limitation:** Can only merge, never split (no "undo")

Thank You!

Contact Information

- Email: ekrem.cetinkaya@yildiz.edu.tr
- Office Hours: Tuesday 14:00-16:00 - Room F-B21
- Course Repo: [GitHub Link](#)